

Cluster

About
Nodes
Node Labels
Applications
NEW
NEW_SAVING
SUBMITTED
ACCEPTED
RUNNING
FINISHED
FAILED
KILLED
Scheduler

Tools

hadoop-master:8088/cluster/app/application_1624479170912_0002

70%

...

...

...

Application Overview

Application application_1624479170912_0002

User: [hadoop](#)

Name: [org.apache.spark.examples.SparkPi](#)

Application Type: [SPARK](#)

Application Tags:

Application Priority: [0 \(Higher Integer value indicates higher priority\)](#)

YarnApplicationState: [FINISHED](#)

Queue: [default](#)

FinalStatus Reported by AM: [SUCCEEDED](#)

Started: [Thu Jun 24 04:23:02 +0800 2021](#)

Launched: [Thu Jun 24 04:23:04 +0800 2021](#)

Finished: [Thu Jun 24 04:23:50 +0800 2021](#)

Elapsed: [48sec](#)

Tracking URL: [History](#)

Log Aggregation Status: [DISABLED](#)

Application Timeout (Remaining Time): [Unlimited](#)

Diagnostics:

Unmanaged Application: [false](#)

Application Node Label expression: [<Not set>](#)

AM container Node Label expression: [<DEFAULT_PARTITION>](#)

Application Metrics

Total Resource Preempted: [<memory:0, vCores:0>](#)

Total Number of Non-AM Containers Preempted: [0](#)

Total Number of AM Containers Preempted: [0](#)

Resource Preempted from Current Attempt: [<memory:0, vCores:0>](#)

Number of Non-AM Containers Preempted from Current Attempt: [0](#)

Aggregate Resource Allocation: [304846 MB-seconds, 148 vcore-seconds](#)

Aggregate Preempted Resource Allocation: [0 MB-seconds, 0 vcore-seconds](#)


Showing 1 to 1 of 1 entries

Search:

Attempt ID	Started	Node	Logs	Nodes blacklisted by the app	Nodes blacklisted by the system
appattempt_1624479170912_0002_000001	Thu Jun 24 04:23:02 +0800 2021	http://hadoop-master:8042	Logs	0	0

Showing 1 to 1 of 1 entries

First Previous 1 Next Last



Nodes of the cluster (unhealthy)

Logged in as: dr.who

Cluster

About
Nodes
Node Labels
Applications
NEW
NEW_SAVING
SUBMITTED
ACCEPTED
RUNNING
FINISHED
FAILED
KILLED
Scheduler

Tools

Cluster Metrics

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Memory Used	Memory Total	Memory Reserved	VCores Used	VCores Total	VCores Reserved
0	0	0	0	0	0 B	24 GB	0 B	0	24	0

Cluster Nodes Metrics

Active Nodes	Decommissioning Nodes	Decommissioned Nodes	Lost Nodes	Unhealthy Nodes	Rebooted Nodes	Shutdown Nodes
3	0	0	0	1	0	0

Scheduler Metrics

Scheduler Type	Scheduling Resource Type	Minimum Allocation	Maximum Allocation	Maximum Cluster Application Priority
Capacity Scheduler	[memory-mb (unit=Mi), vcores]	<memory:1024, vCores:1>	<memory:8192, vCores:4>	0

Show 20 entries

Search:

Node Labels	Rack	Node State	Node Address	Node HTTP Address	Last health-update	Health-report	Containers	Allocation Tags	Mem Used	Mem Avail	VCores Used	VCores Avail	Version
/default-rack		UNHEALTHY	hadoop-slave2:33111	hadoop-slave2:8042	Wed Jun 23 20:20:51 +0800 2021	1/1 local-dirs usable space is below configured utilization percentage/no more usable space [/tmp/hadoop-hadoop/nm-local-dir : used space above threshold of 90.0%]; 1/1 log-dirs usable space is below configured utilization percentage/no more usable space [/usr/local/hadoop /logs/userlogs : used space above threshold of 90.0%]	0		0 B	0 B	0	0	3.2.2

Showing 1 to 1 of 1 entries

First Previous 1 Next Last

Ex.2 Simple Drill Queries

We first generate a large file with students and their grades.

```
import random

firstnames = []
for line in open("firstnames.txt"):
    firstnames.append(line.strip())

lastnames = []
for line in open("lastnames.txt"):
    lastnames.append(line.strip())

for fileNum in range(0, 1):
    fileName = "mapper" + str(fileNum) + ".csv"
    names = []
    IDs = []
    f = open(fileName, "w")
    for i in range(0, 100000): # total number of students is 100000
        ID = ""
        for x in range(0, 10):
            digit = random.randint(0, 9)
            ID = ID + (str(digit))

        names.append(random.choice(firstnames) + " " + random.choice(lastnames))
        IDs.append(ID)

    for j in range(0, 200000000): # total number of lines is 200000000
        grade = random.randint(0, 100)
        index = random.randint(0, 99999)
        result = names[index] + "," + IDs[index] + "," + str(grade) + "\n"
        f.write(result)
```

Then we get a file `mapper0.csv` , which exceeds 5GB.

```
-rw-rw-r-- 1 hadoop hadoop 5741434141 Jun 28 00:07 mapper0.csv
```

By running

```
hdfs dfs -mkdir l4
hdfs dfs -put ./mapper0.csv l4
```

we put the file to hdfs.

We can check <http://10.119.6.238:8047/storage> for storage plugin configuration

2. 1)student who had the lowest grade:

```
apache drill> select name, min(score) as min_score from (select columns[0] as name, INTEGER columns[2] as
score from hdfs.`/user/pgroup1/l4/mapper0.csv`) group by name order by min_score asc limit 1;
+-----+-----+
|      name      | min_score |
+-----+-----+
| Dwana Harradon |    0      |
+-----+-----+
1 row selected (31.429 seconds)
```

2)student who had the highest average grade:

```
apache drill> select name, AVG(score) as avg_score from (select columns[0] as name, cast(columns[2] as int)
as score from hdfs.`/user/pgroup1/l4/mapper0.csv` ) group by name order by avg_score desc limit 1;
+-----+-----+
|      name      | avg_score |
+-----+-----+
| Domonique Mess | 53.2076209086468 |
+-----+-----+
1 row selected (48.804 seconds)
```

3. median over all the scores:

```
apache drill> SELECT COUNT(*) FROM hdfs.`/user/pgroup1/l4/mapper0.csv`;
+-----+
|  EXPR$0  |
+-----+
| 200000000 |
+-----+
1 row selected (22.093 seconds)
apache drill> SELECT AVG(score)
2..semicolon> FROM (SELECT cast(columns[2] as int) as score
3.....)> FROM hdfs.`/user/pgroup1/l4/mapper0.csv`
4.....)> ORDER BY score
5.....)> LIMIT 2 -- odd 1, even 2
6.....)> OFFSET 99999999)
7..semicolon> ;
+-----+
|  EXPR$0  |
+-----+
| 50.0     |
+-----+
1 row selected (85.175 seconds)
```

Ex.3 Simple Spark

(The contents below is identical to the [ex3.ipynb](#) file.)

```
"""
We first run `hdfs dfs -put ./grade.csv /` to put the local file "grade.csv" into hdfs.
"""

from pyspark import SparkConf, SparkContext
sc = SparkContext.getOrCreate(SparkConf().setMaster("local[*]")) # create an instance
```

```
# create an RDD from file
content = sc.textFile("grade.csv")

# .collect() serialize the RDD into a <list>
content.collect()

def f(x):
    a = x.split(',')
    return (a[1], int(a[2]))

mp = content.map(f)
test = mp.reduceByKey(lambda x, y: y if y>x else x)
test.collect()
"""
Check ex3_result for the result.
We don't show the result here because it's too long.
"""
```

```
"""
flatmap() function return a serialized result without Key-Value pair structure.
"""
fm = content.flatMap(f)
fm.collect()
```